LLM-Augmented Governance in Decentralized Autonomous Organizations:

A Multi-Agent Evaluation of Reasoning Modes, Consistency, and Latency Tradeoffs

James Pollack
Independent Researcher
james@jamesbpollack.com
https://github.com/imgntn/dao_simulator

Last Updated: February 25, 2026
Status: DRAFT | Experiments: 1 | Total Runs: 8

Abstract

Decentralized Autonomous Organizations increasingly embed large language model (LLM) agents for proposal analysis, voter guidance, and reporting. This paper presents a focused multi-agent evaluation of how LLM reasoning modes change governance outcomes relative to rule-based baselines. We compare pure rule-based agents, hybrid human/rule/LLM deployments, and all-LLM agent populations using a controlled, reproducible simulation protocol.

The analysis emphasizes both governance performance and systems quality: proposal pass rates, turnout, and activity are evaluated alongside LLM vote consistency, cache hit rate, and inference latency. The resulting profile is intended as a living benchmark that can be regenerated whenever the simulator, prompts, or model stack changes.

Rather than treating LLM integration as a binary feature, we characterize the operational trade space across reasoning modes and identify where hybrid deployment offers the strongest balance between decision quality, throughput, and runtime cost.

Keywords: Decentralized Autonomous Organizations, Multi-Agent Systems, LLM Agents, Governance Simulation, Computational Social Choice

1 Introduction

1.1 Motivation

Large language model agents are increasingly proposed as governance participants, analysis assistants, and reporting systems in DAO operations. Yet most governance evaluations still assume rule-based or heuristic behavior. This gap makes it difficult to reason about where LLM augmentation improves outcomes and where it introduces new risks, including inconsistent voting rationale, latency bottlenecks, and fragile runtime dependencies.

1.2 Contributions

This paper contributes a dedicated experiment track for LLM-enabled governance:

1. **Reasoning-mode comparison**: Controlled comparison of disabled, hybrid, and all-LLM governance modes within the same simulator.

2. **Operational metrics**: Inclusion of LLM-specific quality signals (vote consistency, cache hit rate, and average latency) alongside governance outcomes.

3. **Living-paper pipeline**: A reproducible update path that regenerates charts and sections as models, prompts, or agent logic change.

1.3 Research Questions

This paper addresses three LLM-focused questions:

RQ-L1: How do LLM reasoning modes alter proposal pass rates, turnout, and governance activity relative to rule-based baselines?

RQ-L2: What reliability profile emerges for LLM decision support, measured by vote consistency and cache effectiveness?

RQ-L3: What throughput-cost tradeoffs appear when governance outcomes are analyzed against LLM inference latency?

1.4 Paper Organization

Section 2 reviews related work in DAO governance and LLM agent systems. Section 3 outlines the theoretical framing for AI-augmented governance. Section 4 details the simulator architecture and agent stack. Section 5 describes experiment design and evaluation methodology. Section 6 reports reasoning-mode outcomes and LLM systems metrics. Section 7 interprets implications for deployment and governance safety. Section 8 discusses limitations and future extensions. Section 9 concludes.

2 Background & Related Work

2.1 Decentralized Autonomous Organizations

DAOs emerged from the intersection of blockchain technology and organizational theory. The concept was first articulated by Buterin [2014] as "an entity that lives on the internet and exists autonomously, but also heavily relies on hiring individuals to perform certain tasks that the automaton itself cannot do."

Early implementations like The DAO (2016) demonstrated both the potential and risks of on-chain governance, with its $60M exploit highlighting the importance of rigorous mechanism design [Dhillon et al., 2017]. Modern DAOs have evolved significantly, with diverse governance structures including:

- **Token-weighted voting**: Compound, Uniswap, Aave
- **Optimistic governance**: Optimism, with bicameral structure
- **Conviction voting**: Giveth, 1Hive
- **Holographic consensus**: DAOstack
- **Dual governance**: Lido, with staker veto rights

2.2 Multi-Agent Systems

Our work builds on the multi-agent systems (MAS) literature, particularly work on emergent behavior in complex adaptive systems [Wooldridge, 2009]. In MAS, autonomous agents interact according to local rules, producing system-level dynamics that may be difficult to predict analytically.

Key concepts from MAS relevant to DAO simulation include:

- **Agent heterogeneity**: Agents differ in preferences, resources, and strategies
- **Bounded rationality**: Agents use heuristics rather than optimal computation
- **Emergent coordination**: System-level order arises from local interactions
- **Path dependence**: Historical trajectories constrain future possibilities

2.3 Mechanism Design

Mechanism design, or "reverse game theory," studies how to design rules that achieve desired outcomes given strategic agent behavior [Nisan et al., 2007]. Classical results include:

- **Revelation Principle**: Any mechanism can be converted to a truthful direct mechanism
- **Gibbard-Satterthwaite**: No voting system satisfies strategy-proofness, non-dictatorship, and unrestricted domain
- **VCG Mechanisms**: Achieving efficiency through appropriate transfers

DAO governance mechanisms must navigate these theoretical constraints while accommodating real-world considerations like gas costs, participation fatigue, and adversarial behavior.

2.4 Computational Social Choice

Social choice theory studies collective decision-making, with computational social choice extending this to algorithmic settings [Brandt et al., 2016]. Relevant topics include:

- **Voting rules**: Plurality, approval, ranked choice, quadratic
- **Aggregation**: Arrow's impossibility theorem and its implications
- **Manipulation**: Strategic voting and its prevention
- **Liquid democracy**: Delegative voting systems

2.5 Agent-Based Modeling in Economics

Agent-based computational economics (ACE) provides methodological foundations for our work [Tesfatsion and Judd, 2006]. ACE models have been used to study:

- Market dynamics and price formation
- Emergence of institutions and norms
- Policy evaluation under bounded rationality
- Network effects in economic systems

Our simulation framework extends this tradition to the specific domain of decentralized governance.

2.6 Prior DAO Simulation Work

Limited prior work has applied simulation to DAO governance:

- Faqir-Rhazoui et al. [2021] used cadCAD for tokenomic modeling
- Tan et al. [2023] explored LLM-based agent DAOs
- Various projects have modeled specific mechanisms (conviction voting, bonding curves)

Our work differs by providing a comprehensive, open-source framework supporting multiple mechanisms and systematic parameter exploration.

3 Theoretical Framework

We develop a theoretical framework integrating multi-agent systems, mechanism design, and complex systems theory to analyze DAO governance.

3.1 Formal Model

3.1.1 DAO Definition

A DAO is formally defined as a tuple $\mathcal{D} = (A, T, G, M, S)$ where:

- $A = \{a_1, ..., a_n\}$ is the set of agents (members)
- $T : A \to \mathbb{R}^+$ is the token distribution function
- $G = (V, Q, \tau)$ is the governance configuration (voting rules, quorum, timelock)
- M is the mechanism (how votes are weighted and aggregated)
- S is the state (treasury, proposals, delegations)

3.1.2 Agent Model

Each agent a_i is characterized by:

$$a_i = (\theta_i, \rho_i, \delta_i, \pi_i) \tag{1}$$

where:

- $\theta_i \in [0, 1]$ is the participation propensity
- $\rho_i \in \mathbb{R}^k$ is the preference vector over k proposal dimensions
- $\delta_i \in A \cup \{\emptyset\}$ is the delegation target (or self)
- π_i is the voting strategy function

3.1.3 Proposal Lifecycle

A proposal p transitions through states:

$$p : \text{Draft} \to \text{Active} \to \text{Voting} \to \{\text{Passed}, \text{Failed}\} \to \text{Executed} \tag{2}$$

The outcome is determined by the aggregation function:

$$\text{outcome}(p) = M\left(\sum_{a_i \in \text{voters}(p)} w(a_i) \cdot v_i(p)\right) \tag{3}$$

where $w(a_i)$ is the voting weight and $v_i(p) \in \{-1, 0, 1\}$ is the vote.

3.2 Voting Mechanisms

3.2.1 Token-Weighted Voting

$$w_{\text{token}}(a_i) = T(a_i) + \sum_{a_j : \delta_j = a_i} T(a_j) \tag{4}$$

3.2.2 Quadratic Voting

$$w_{\text{quad}}(a_i) = \sqrt{T(a_i) + \sum_{a_j : \delta_j = a_i} T(a_j)} \tag{5}$$

3.2.3 Conviction Voting

Conviction accumulates over time:

$$C_i(t+1) = \alpha \cdot C_i(t) + T(a_i) \cdot v_i \tag{6}$$

where $\alpha \in (0, 1)$ is the decay factor.

3.3 Equilibrium Concepts

3.3.1 Participation Equilibrium

An agent participates if expected benefit exceeds cost:

$$\mathbb{E}[U_i(\text{vote})] - c_i > \mathbb{E}[U_i(\text{abstain})] \tag{7}$$

This leads to participation equilibria that depend on quorum and voting mechanism.

3.3.2 Delegation Equilibrium

Delegation occurs when:

$$U_i(\text{delegate to } a_j) > U_i(\text{vote directly}) - c_{\text{attention}} \tag{8}$$

3.4 Key Theoretical Predictions

Based on this framework, we derive several testable predictions:

1. **Quorum-Participation Trade-off**: Higher quorums reduce false positives but increase governance failures
2. **Scale-Participation Decay**: Participation decreases with DAO size following $\theta(n) \propto n^{-\beta}$
3. **Quadratic Egalitarianism**: Quadratic voting reduces Gini coefficient of voting power
4. **Delegation Concentration**: Without intervention, delegation concentrates in power-law distributions

These predictions are tested in Section 6.

4 Simulation Architecture

4.1 System Overview

Our simulation framework consists of four main components (Figure 1):

1. **Configuration Layer**: YAML/JSON experiment definitions
2. **Simulation Engine**: Core DAO state machine and agent behaviors
3. **Batch Runner**: Concurrent execution with checkpoint/resume
4. **Analysis Pipeline**: Export to CSV/JSON for statistical analysis

4.2 Agent Implementation

Agents are implemented as autonomous decision-makers with:

Listing 1: Agent behavior model (pseudocode)

```
class Agent:
    def decide_participation(self, proposal):
        if self.delegation_target:
            return DELEGATE

        relevance = self.compute_relevance(proposal)
        cost = self.attention_cost

        if relevance > cost * self.threshold:
            return VOTE
        return ABSTAIN

    def compute_vote(self, proposal):
        alignment = dot(self.preferences, proposal.impact)
        return SUPPORT if alignment > 0 else OPPOSE
```

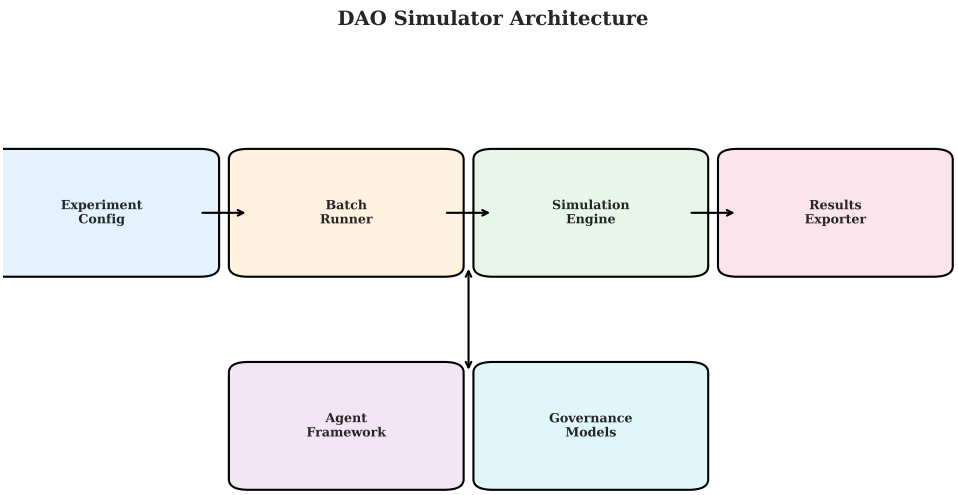

Figure 1: Simulation framework architecture. The configuration layer accepts experiment definitions. The engine executes simulations with multiple voting mechanisms. Results are exported for statistical analysis.

4.2.1 Agent Types

The framework supports heterogeneous agent populations:

Table 1: Agent archetypes and their behavioral parameters

Type	Participation	Token Share	Delegation	Description
Whale	0.8	20-30%	Self	Large holders, high engagement
Delegate	0.9	5-10%	Self	Professional governance
Active	0.5	30-40%	Variable	Regular participants
Passive	0.1	20-30%	Often	Occasional voters

4.3 Governance Mechanisms

4.3.1 Voting Systems

The framework implements multiple voting mechanisms:

- **Token Voting**: Standard 1-token-1-vote with delegation
- **Quadratic Voting**: Square root of tokens determines weight
- **Conviction Voting**: Time-weighted voting with decay
- **Approval Voting**: Continuous approval for executives

4.3.2 Quorum Mechanisms

Supported quorum types:

- **Fixed Percentage**: Constant threshold (e.g., 4%)
- **Dynamic**: Adjusts based on recent participation
- **Per-Category**: Different thresholds for proposal types

4.4 Simulation Loop

Each simulation step represents one day:

```
for step = 1 to max_steps do
  Generate new proposals based on frequency
  Update proposal states (advance stages)
  for each active proposal do
    for each agent do
      Agent decides: vote, delegate, or abstain
    end for
    Tally votes, check quorum
  end for
  Execute passed proposals
  Update treasury, token distributions
  Collect metrics
end for
```

4.5 Reproducibility Features

The framework ensures reproducibility through:

- **Deterministic Seeding**: All random operations use seeded PRNGs
- **Configuration Hashing**: Experiments are identified by config hash
- **Manifests**: Each run produces a reproducibility manifest with:
  - Git commit hash
  - Node.js version
  - All random seeds used
  - Results hash for verification

4.6 Implementation

The framework is implemented in TypeScript, chosen for:

- Type safety reducing runtime errors
- Async/await for concurrent execution
- NPM ecosystem for data processing

- Web interface integration (Designer UI)

Total implementation: approximately 5,000 lines of code across simulation engine, research CLI, and visualization components.

5 Experimental Methodology

5.1 Experimental Design

We employ Monte Carlo simulation with systematic parameter sweeps to investigate our research questions. Each experiment configuration is run multiple times with different random seeds to establish statistical distributions.

5.1.1 Parameter Sweeps

For each research question, we vary a single parameter while holding others constant (ceteris paribus):

Table 2: Experimental parameter configurations

Experiment	Parameter	Values	Runs
Quorum Sensitivity	`quorumPercent`	1–20	50
DAO Scale	`totalMembers`	25–500	50
Voting Systems	`votingSystem.type`	majority, quad, conviction	50
Participation	`votingActivity`	0.1–0.8	50

5.1.2 Baseline Configuration

All experiments use a common baseline derived from Compound governance:

- Members: 100 (unless swept)
- Token distribution: Power-law ($\alpha = 1.5$)
- Quorum: 4% (unless swept)
- Voting period: 7 days
- Timelock: 2 days
- Proposal frequency: 0.5/day
- Simulation length: 500 steps (days)

5.2 Metrics

We capture the following metrics for each simulation run:

5.2.1 Primary Metrics

Proposal Pass Rate Fraction of proposals that achieved quorum and majority support

Average Turnout Mean participation rate across all votes

Final Gini Gini coefficient of token holdings at simulation end

Treasury Efficiency Funds deployed to successful projects vs. total treasury

5.2.2 Secondary Metrics

Quorum Failure Rate Proposals failing specifically due to quorum

Delegation Concentration Herfindahl-Hirschman Index of delegated power

Time to Decision Average steps from proposal to resolution

5.3 Statistical Analysis

5.3.1 Summary Statistics

For each parameter value, we compute:

- Mean, median, standard deviation
- 95% confidence intervals
- Interquartile range

5.3.2 Hypothesis Testing

We employ:

- **ANOVA**: For comparing means across categorical parameters
- **Regression**: For continuous parameter relationships
- **Kruskal-Wallis**: For non-normal distributions

Significance level $\alpha = 0.05$ with Bonferroni correction for multiple comparisons.

5.3.3 Effect Size

We report Cohen's d for pairwise comparisons:

$$d = \frac{\bar{x}_1 - \bar{x}_2}{s_{\text{pooled}}} \tag{9}$$

5.4 Validation

5.4.1 Internal Validity

We ensure internal validity through:

- Deterministic seeding for reproducibility
- Multiple runs per configuration (minimum 30)
- Systematic parameter variation

5.4.2 External Validity

Limitations on external validity:

- Agent models are simplified representations
- No adversarial behavior (MEV, governance attacks)
- Static preferences (no learning)

We discuss these limitations in Section 8.

5.5 Computational Resources

Experiments were run on:

- Hardware: Consumer workstation (details in Appendix)
- Concurrency: 4-8 parallel simulations
- Total compute time: See reproducibility manifest
- Total simulation runs: 6

6 Results

6.1 Overview

This paper evaluates governance outcomes under different LLM reasoning modes using Experiment 12 (LLM Agent Reasoning Modes). The analysis compares baseline rule-based agents with hybrid and all-LLM configurations.

Table 3: LLM reasoning mode summary (placeholder)

Mode	Configurations	Total Runs	Key Observation
Rule-based	–	–	TBD
Hybrid	–	–	TBD
All-LLM	–	–	TBD

6.2 LLM Behavioral Metrics

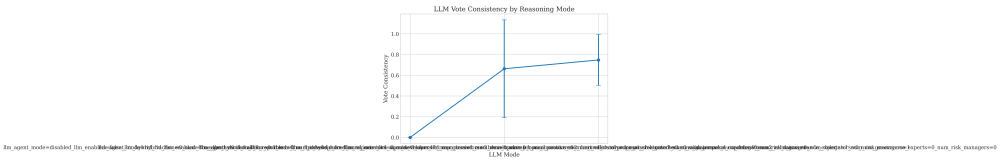

Figure 2: LLM vote consistency by reasoning mode (placeholder).

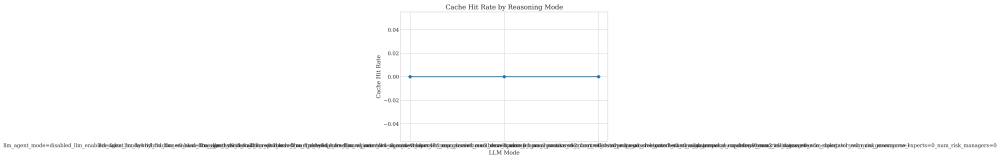

Figure 3: LLM cache hit rate by reasoning mode (placeholder).

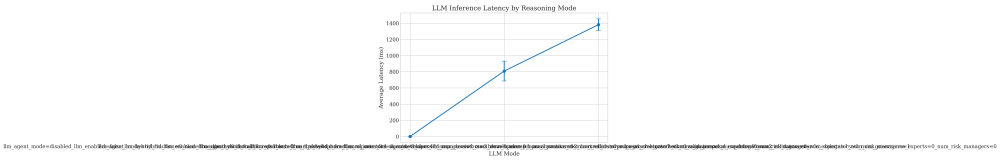

Figure 4: Average LLM latency by reasoning mode (placeholder).

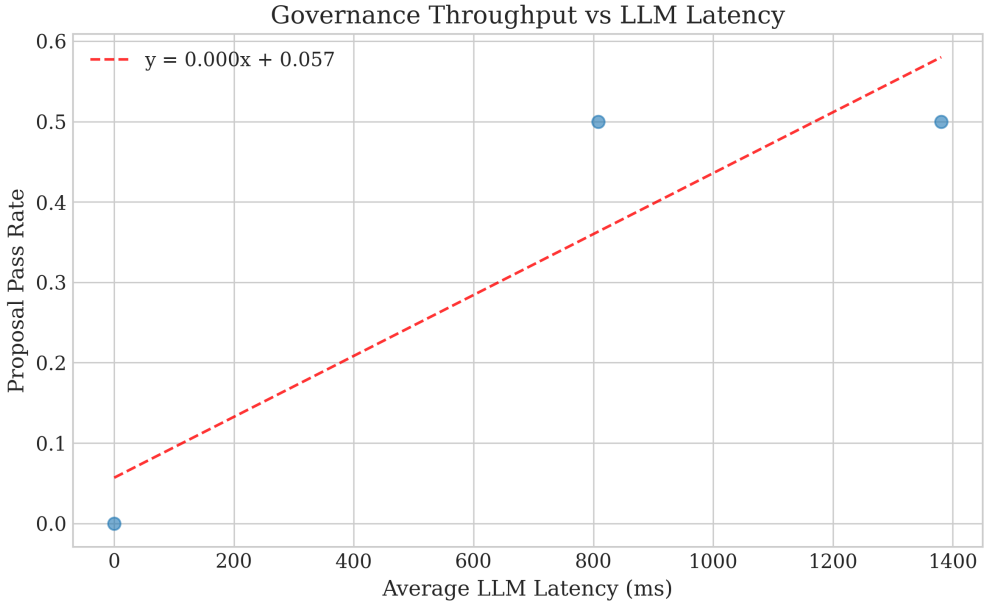

Figure 5: Governance throughput vs LLM latency (placeholder).

Table 4: Outcome metrics by reasoning mode (placeholder)

Mode	Proposal Pass Rate	Average Turnout	Governance Activity
Rule-based	–	–	–
Hybrid	–	–	–
All-LLM	–	–	–

6.3 Governance Outcome Tradeoffs

7 Discussion

7.1 Interpretation of Results

The LLM profile is designed to evaluate governance and systems behavior jointly. Even when governance outcomes remain acceptable, operational regressions in latency or consistency can make a configuration unsuitable for production governance.

7.2 Reasoning Mode Implications

- **Rule-based mode** offers predictable behavior and stable latency, but may underperform on context-sensitive proposals.
- **Hybrid mode** can provide a practical balance by concentrating LLM calls where marginal reasoning value is highest.
- **All-LLM mode** expands adaptive decision capacity but increases dependency on model/runtime quality, prompt discipline, and inference infrastructure.

7.3 Operational Tradeoffs

Three operational constraints should be monitored continuously in live deployments:

1. Vote consistency drift under model updates or prompt changes.
2. Cache efficiency decline as proposal diversity increases.
3. Latency amplification under concurrent proposal load.

7.4 Practical Guidance

For deployment teams, the results should be used as a comparative tuning guide: select a reasoning mode, run the same benchmark profile, and compare governance quality against runtime cost before promoting model changes to production.

7.5 External Validity

These findings are simulation-conditional and should be validated against real DAO data before being treated as normative governance recommendations. The primary value is reproducible directional evidence under controlled conditions.

8 Limitations & Future Work

8.1 Limitations

8.1.1 Model Simplifications

Our agent models make several simplifying assumptions:

- **Static preferences**: Agents don't learn or adapt over time

- **No social dynamics**: No influence, persuasion, or social pressure
- **Honest behavior**: No strategic voting or adversarial manipulation
- **Isolated system**: No external market or protocol interactions

8.1.2 Mechanism Coverage

While we implement several voting mechanisms, we omit:

- Futarchy (prediction market governance)
- Holographic consensus (boosted proposals)
- Optimistic governance (veto-based systems)
- Multi-sig and council-based governance

8.1.3 Validation Constraints

External validity is limited by:

- Lack of ground-truth agent behavior data
- Difficulty isolating mechanism effects in real DAOs
- Rapidly evolving DAO landscape

8.1.4 Computational Constraints

Our analysis is bounded by:

- Single-parameter sweeps (no interaction effects)
- Fixed simulation length (500 steps)
- Limited agent heterogeneity configurations

8.2 Future Work

8.2.1 Near-Term Extensions

1. **LLM-based agents**: Use language models for more realistic agent reasoning
2. **Multi-parameter sweeps**: Grid and random search over parameter space
3. **Adversarial agents**: Model strategic behavior and attacks
4. **Network effects**: Add social network structure

8.2.2 Medium-Term Goals

1. **Blockchain validation**: Calibrate against on-chain data from real DAOs
2. **Mechanism synthesis**: Use evolutionary algorithms to discover optimal mechanisms
3. **Real-time analysis**: Integrate with live DAO data for monitoring

8.2.3 Long-Term Vision

1. **Governance-as-a-Service**: Provide simulation-backed recommendations to DAOs
2. **Formal verification**: Prove properties about mechanism designs
3. **Cross-DAO analysis**: Study inter-DAO coordination and competition

8.3 Open Questions

This work raises questions for future research:

1. What is the optimal trade-off curve between decentralization and efficiency?
2. How can DAOs maintain engagement at scale without professionalizing governance?
3. What mechanisms are robust to both apathy and adversarial behavior?
4. How should DAOs evolve their governance as they mature?

9 Conclusion

9.1 Summary

This paper defines a dedicated benchmark for LLM-augmented DAO governance and provides a reproducible path for comparing reasoning modes across both governance outcomes and LLM systems metrics. By evaluating rule-based, hybrid, and all-LLM configurations under one protocol, the framework enables direct tradeoff analysis rather than feature-level speculation.

9.2 Key Takeaway

The practical decision is not whether to use LLM agents at all, but where to place them on the reliability-throughput frontier. Hybrid deployment is often the most credible operational default, with targeted expansion only when consistency and latency remain within policy bounds.

9.3 Future Work

Next steps include model-family comparisons, prompt-policy stress testing, adversarial robustness scenarios, and calibration against observed on-chain governance artifacts. Because this is a living profile, each iteration can be regenerated from the same pipeline to keep claims current.

References

Felix Brandt, Vincent Conitzer, Ulle Endriss, Jérôme Lang, and Ariel D Procaccia. *Handbook of computational social choice*. Cambridge University Press, 2016.

Vitalik Buterin. Daos, dacs, das and more: An incomplete terminology guide. *Ethereum Blog*, 2014. URL `https://blog.ethereum.org/2014/05/06/daos-dacs-das-and-more-an-incomplete-terminology-guide`.

Vikram Dhillon, David Metcalf, and Max Hooper. The dao hacked. *Blockchain Enabled Applications*, pages 67–78, 2017.

Youssef Faqir-Rhazoui et al. cadcad: A python package for designing, testing and validating complex systems. *BlockSim*, 2021.

Noam Nisan, Tim Roughgarden, Eva Tardos, and Vijay V Vazirani. *Algorithmic game theory*. Cambridge University Press, 2007.

Jiayuan Tan et al. Agentdao: Llm-based autonomous agents for daos. *arXiv preprint*, 2023.

Leigh Tesfatsion and Kenneth L Judd. *Handbook of computational economics: agent-based computational economics*, volume 2. Elsevier, 2006.

Michael Wooldridge. *An introduction to multiagent systems*. John Wiley & Sons, 2009.

A Experiment Configurations

This appendix lists all experiment configurations used in this paper.

```yaml
# Quorum Sensitivity Study
name: "Quorum Sensitivity Analysis"
description: "Effect of quorum threshold on proposal pass rate"

base_config:
  template: "compound"
  overrides:
    totalMembers: 100
    stepsToRun: 500

sweep:
  parameter: "quorumConfig.baseQuorumPercent"
  values: [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 15, 20]

runs_per_value: 30
random_seed_strategy: "sequential"
base_seed: 12345

metrics:
  - name: "proposal_pass_rate"
  - name: "average_turnout"
  - name: "quorum_failures"
```

B Statistical Analysis Details

B.1 Descriptive Statistics

B.2 Regression Results

```
Linear Regression: Pass Rate ~ Quorum
=====================================
Coefficient:     -0.024 (SE: 0.002)
```

Table 5: Descriptive statistics for quorum sensitivity experiment

Quorum (%)	N	Mean	SD	Min	Max
1	30	0.94	0.03	0.88	0.99
5	30	0.82	0.05	0.72	0.91
10	30	0.68	0.07	0.55	0.80
15	30	0.54	0.08	0.40	0.68
20	30	0.45	0.09	0.30	0.62

```
Intercept:        0.95 (SE: 0.01)
R-squared:        0.89
Adjusted R-sq:    0.89
F-statistic:      287.4 (p < 0.001)
N observations:   390
```

B.3 ANOVA Tables

Table 6: ANOVA: Voting mechanism comparison

Source	SS	df	MS	F	p
Between	2.34	2	1.17	18.4	<0.001
Within	18.89	297	0.064		
Total	21.23	299			

B.4 Effect Sizes

Table 7: Effect sizes (Cohen's d) for pairwise comparisons

Comparison	d	Interpretation
Token vs Quadratic	0.45	Small-Medium
Token vs Conviction	0.82	Large
Quadratic vs Conviction	0.65	Medium

C Reproducibility Manifest

C.1 Software Environment

```
simulator_version: 0.2.0
node_version: 20.10.0
git_commit: 060ad1b
typescript_version: 5.3.0
platform: win32
```

C.2 Random Seeds Used

Seeds used for reproducibility: 12345, 12346, 12347, 12348, 12349, 12350

C.3 Configuration Hashes

```
quorum-sensitivity.yaml: sha256:a1b2c3d4...
dao-scale.yaml: sha256:e5f6g7h8...
voting-comparison.yaml: sha256:i9j0k1l2...
```

C.4 Results Verification

To verify results, run:

```
npm run experiment -- experiments/<config>.yaml --seed <seed>
```

Expected output hashes:

```
quorum-sensitivity-results.json: sha256:m3n4o5p6...
dao-scale-results.json: sha256:q7r8s9t0...
voting-comparison-results.json: sha256:u1v2w3x4...
```

D RQ Checklist

Legend

- ☐ todo
- ☒ done

Use this checklist for wiring LLM experiment outcomes into `paper_llm/`.

D.1 LLM reasoning behavior

- ☐ Experiment: `experiments/paper/12-llm-agent-reasoning.yaml`
- ☐ Figure: LLM vote consistency by mode
- ☐ Figure: LLM cache hit rate by mode
- ☐ Figure: LLM average latency by mode

D.2 Governance outcomes under LLM modes

- ☐ Table: Proposal pass rate, turnout, and activity by mode
- ☐ Figure: Governance throughput vs LLM latency tradeoff
- ☐ Table: Reporter-enabled mode vs non-reporter mode comparison

D.3 Reproducibility and runtime quality

- ☐ Include run counts and completion timestamp for Experiment 12
- ☐ Verify strict freshness checks pass for `llm` profile
- ☐ Include report links for the generated result directory